

As mentioned in the previous section, AI is the broader concept that aims to create systems that can perform tasks akin to humans. Machine learning is a way to learn from and make informed predictions and decisions based on data. It can be thought of as yet another form of automation that involves using programs to learn and improve over time without explicit programming. Finally, Data Science, as a

**MODULE 9**

# KUBERNETES (K8s)

Thinking about it broadly, we could divide applications in AI, ML, and data science into two broad categories, namely Predictive AI and Generative AI. Predictive AI aims at predicting and analyzing outcomes (e.g., classification, clustering, regression, object detection, etc.). In contrast, generative AI aims at generating new and original content (e.g., LLMs, RAG<sup>17</sup>, etc.). As such, the algorithms and techniques underpinning predictive and generative AI can vary widely.

## ZERO TO HERO

Master Container Orchestration from Basics to Advanced

### Deploy

Automated Deployment

### Scale

Auto-scaling Magic

### Heal

Self-healing Systems

Workloads

Classification Object Detection

RAGs LLMs



# Learning Journey

---

01

## Introduction & Core Concepts

- What is Kubernetes?
- Problems K8s Solves
- Container Orchestration Features
- Core Components Overview

02

## Architecture & Components

- Master Node vs Worker Node
- K8s Architecture Deep Dive
- Local Setup with Minikube
- YAML Configuration

03

## Practical Deployment

- Kubectl Essential Commands
- Demo Project: Fenrir App
- Namespaces & Ingress
- Data Persistence & Services

04

## Advanced Ecosystem

- Helm Package Manager
- ArgoCD for GitOps
- Istio Service Mesh
- Real DevOps Workflow

# What is Kubernetes?

## Definition

Kubernetes (K8s) is an open-source container orchestration platform that automates the deployment, scaling, and management of containerized applications.

 Docker runs containers →  K8s manages containers

## Key Facts

- ✓ Works with Docker, containerd
- ✓ Used by Google, Netflix, Amazon
- ✓ Manages clusters of machines

## Problems Kubernetes Solves

- ✗ Manual deployment
  - ✗ No auto-scaling
  - ✗ Difficult networking
  - ✗ Downtime in updates
- ✓ Automatic scaling
  - ✓ Self-healing
  - ✓ Load balancing
  - ✓ Zero downtime

## Container Orchestration Features



### Auto-scheduling



### Auto-scaling



### Self-healing



### Load balancing



### Rolling updates



### Config management



### Service discovery



### Network management

# Core Components & Architecture

**Pod**

Smallest deployable unit containing one or more containers.

**ConfigMap**

Stores configuration data.

**Deployment**

Manages stateless apps with auto-scaling.

**Service**

Exposes pods and provides stable networking.

**Secret**

Stores sensitive data like passwords.

**StatefulSet**

Manages stateful apps like databases.

**Ingress**

Manages external HTTP/HTTPS access.

**Volume**

Provides persistent storage.

**Node**

A machine (VM/physical) running pods.

## Master Node (Control Plane)

API Server

Entry point

Scheduler

Assigns pods

Controller Manager

Maintains state

ETCD

Cluster database

## Worker Node

Kubelet

Communicates

Kube Proxy

Networking

Container Runtime

Runs containers

Pods

Your apps

# Local Setup & YAML Configuration

## Minikube

Runs a local K8s cluster on your laptop for development.

```
minikube start
```

## YAML Configuration Structure



### metadata

Name, labels



### spec

Desired state



### status

Auto-generated

## Example Pod YAML Configuration

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
```

# Essential Kubectl Commands

## ☰ Get Resources

```
kubectl get pods
```

List all pods

```
kubectl get nodes
```

List all nodes

```
kubectl get services
```

List all services

## ➕ Create Resources

```
kubectl apply -f app.yaml
```

Create from YAML

```
kubectl create namespace dev
```

Create namespace

## "description" Describe Resources

```
kubectl describe pod mypod
```

Detailed pod info

```
kubectl explain pod
```

Resource schema

## Logs & Debug

```
kubectl logs mypod
```

Pod logs

```
kubectl exec -it mypod -- bash
```

Access pod shell

```
kubectl top nodes
```

Resource usage

## ➖ Delete Resources

```
kubectl delete pod mypod
```

Delete pod

```
kubectl delete -f app.yaml
```

Delete from YAML

## Pro Tips

**Aliases:** k=kubectl

**Shortcuts:** no, pods, svc, services, deploy, deployments

# Demo Project: Deploy Fenrir App

1

**Build**

Create Docker image

2

**Push**

Upload to registry

3

**Deploy**

Create Deployment

4

**Expose**

Create Service

5

**Access**

Via NodePort

## 1. Build & Push Docker Image

```
docker build -t fenrir-app .
```

```
docker tag fenrir-app username/fenrir-app
```

```
docker push username/fenrir-app
```

## 2. Create Deployment YAML

```
apiVersion: apps/v1
kind: Deployment
metadata: name: fenrir-deployment
spec: replicas: 3
selector: matchLabels: app: fenrir
template:
  metadata: labels: app: fenrir
  spec:
    containers:
      - name: fenrir
        image: username/fenrir-app
      ports:
        - containerPort: 3000
```

## 3. Create Service

```
apiVersion: v1
kind: Service
metadata: name: fenrir-service
spec:
  selector: app: fenrir
  type: NodePort
  ports:
```

## 4. Deploy & Verify

```
kubectl apply -f deployment.yaml
```

```
kubectl apply -f service.yaml
```

```
kubectl get pods
```

# Namespaces & Ingress

## What is a Namespace?

Namespaces provide **logical separation** within a single cluster, like folders in a filesystem.

- **default** - Your resources
- **kube-system** - System resources
- **kube-public** - Public info

## Create Namespace

```
kubectl create namespace dev
```

- ✓ Environment isolation
- ✓ Resource limits
- ✓ Security boundaries

## Ingress vs Service

### Service (L4)

TCP/UDP, internal/external

### Ingress (L7)

HTTP/HTTPS, external

### Ingress Controller Required

Ingress resources need an Ingress Controller (like nginx or traefik) to process the routing rules.

## Routing Example

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: app-ingress
spec:
  rules:
    - host: example.com
      http:
        paths:
          - path: /api
            backend:
              service:
                name: backend
                port:
                  number: 80
          - path: /ui
            backend:
              service:
                name: frontend
                port:
                  number: 80
```

# Data Persistence & Services



## Persistent Volume

Actual storage



## PVC

Storage request



## Storage Class

Dynamic provisioning

### PVC Example

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
spec:
  accessModes: ["ReadWriteOnce"]
  resources:
    requests:
      storage: 5Gi
```

### Using in Pod

```
spec:
  containers:
    - name: app
      image: myapp
      volumeMounts:
        - name: storage
          mountPath: /data
      volumes:
        - name: storage
          persistentVolumeClaim:
            claimName: myclaim
```

### Service Types



#### ClusterIP

Internal only



#### NodePort

External (Dev)



#### LoadBalancer

Cloud LB



#### Headless

Direct Pod IP

# Helm Package Manager & ArgoCD



The [package manager](#) for Kubernetes.

## Helm Chart

A bundle of K8s resources.

- ✓ Reusability
- ✓ Easy upgrades
- ✓ Configurable values

## Helm Chart Structure

```
chart/
  └── templates/
    ├── deployment.yaml
    └── service.yaml
  └── values.yaml
└── Chart.yaml
```



GitOps continuous delivery tool.

## Core Principle

Git is the source of truth.

- ✓ Auto-sync cluster
- ✓ Continuous deployment
- ✓ Drift detection

## Workflow

1. Push to Git
2. ArgoCD detects change
3. Auto-sync to cluster

## PRODUCTION WORKFLOW

# Istio & Real DevOps K8s Workflow

## Istio Service Mesh

A **service mesh** for traffic management, security, and observability.



### Traffic Mgmt

Routing, load balancing



### Security

mTLS, policies



### Observability

Metrics, tracing

## Key Benefits

**Canary Deployments:** Gradual rollout

**Circuit Breakers:** Fail fast & recover

**Zero-Trust Network:** mTLS encryption

## Real DevOps K8s Workflow



Code Push  
Git



Build Image  
CI/CD



Push Registry  
Docker Hub



Detect Change  
ArgoCD



Deploy App  
K8s



Expose Service  
Ingress



Monitor  
Prometheus



Verify  
Tests

# Mini Project & Interview Prep

## Mini Project: Flask + MongoDB

Deploy a complete stack: Flask (frontend) + MongoDB (database).

Deployment (Flask)

StatefulSet (Mongo)

Service

ConfigMap

PVC

Ingress

## Project Commands

```
# Deploy all  
kubectl apply -f .  
  
# Check status  
kubectl get all  
  
# Access via Ingress  
curl http://myapp.local
```

## Top Interview Questions

### Q1: What is a Pod?

Smallest deployable unit. Can contain one or more containers.

### Q2: Deployment vs StatefulSet?

Deployment: Stateless apps. StatefulSet: Stateful apps like DBs.

### Q3: What is ETCD?

Distributed key-value store for cluster state.

## More Key Questions

### Q4: What is Ingress?

Manages external HTTP/HTTPS access.

### Q5: What is Helm?

K8s package manager for resource bundles.



# Your K8s Journey Continues

You now have the foundation to orchestrate containers like a pro.

**Remember:** Practice is key to mastery.



## Keep Experimenting

Build, break, and rebuild. True learning comes from doing.



## Join the Community

Connect, share, and grow with fellow K8s enthusiasts.



## Go Build Amazing

The cloud-native world is yours to architect.

🚀 Happy Orchestrating! 🚀